

Internet Research Observatory

W205 Spring 2017 - Final Presentation

Authors: Jason Hunsberger, Victoria Baker, Dominic
Delmolino

4,294,967,296

IPv4 addresses on the public Internet

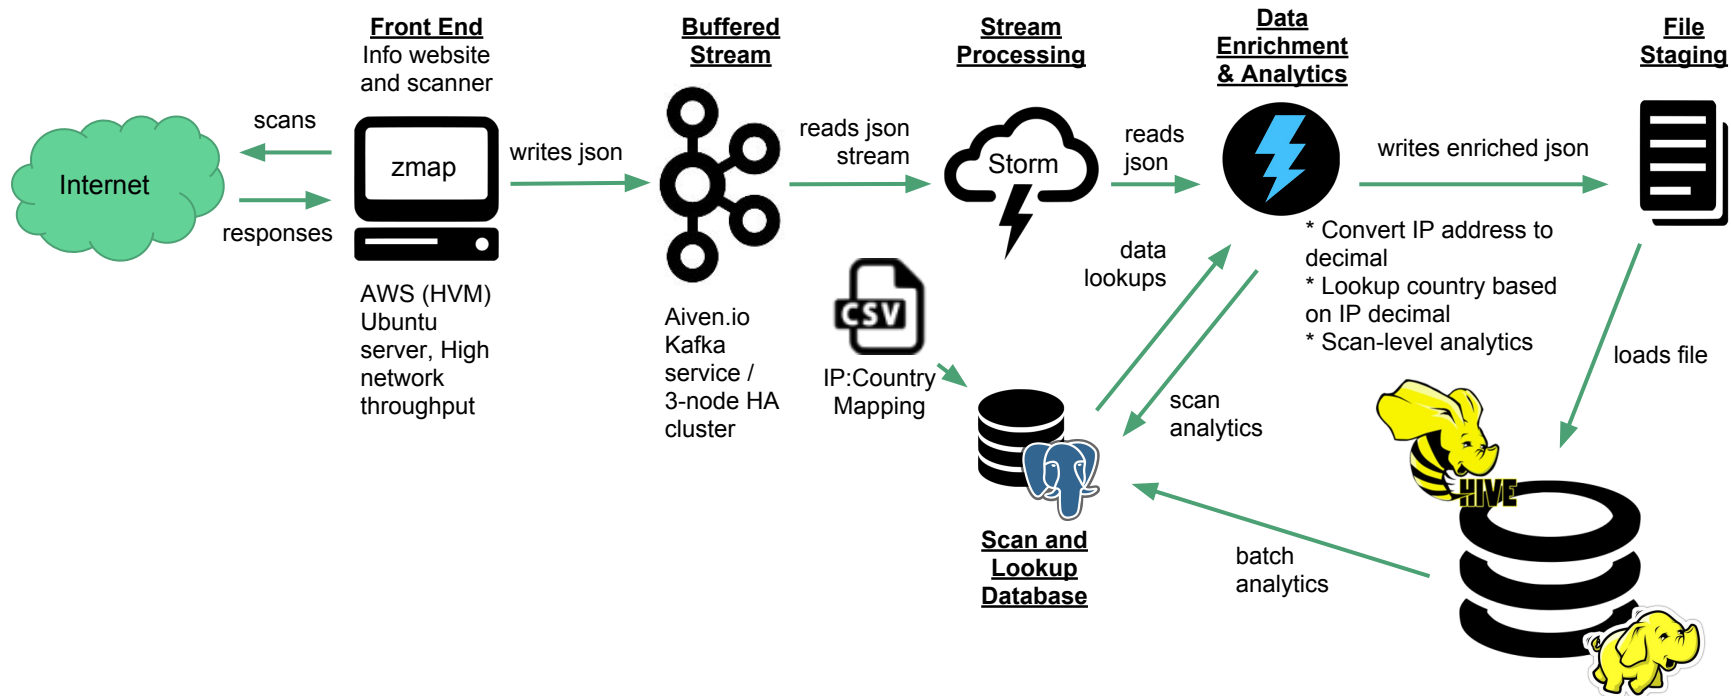
How many of them are actively in use?

What is the geographical distribution of active addresses?

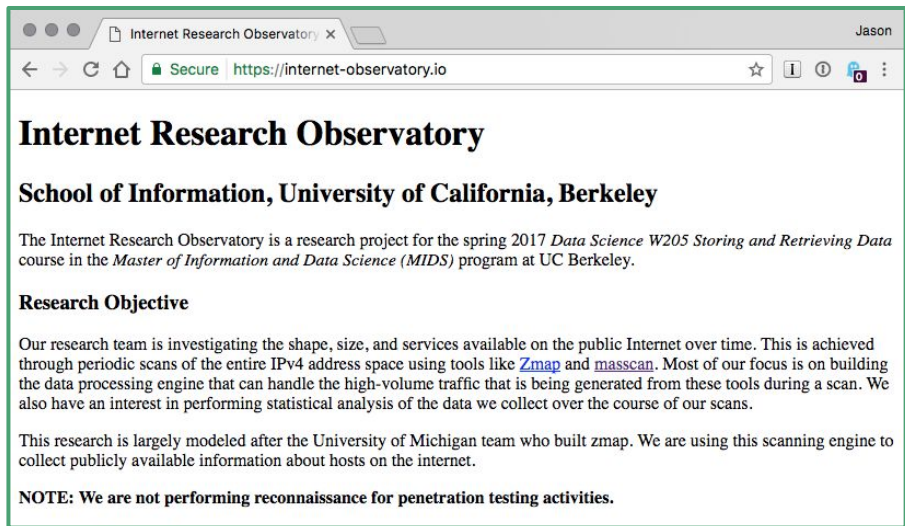
Are there common vulnerabilities across addresses?

***Can we scan all of them in a reasonable amount of time,
record the results and answer these questions?***

Solution Architecture



Scanning the Internet with Zmap



Front End Server

AWS EC2 Ubuntu 16.04 (HVM) w/high-network bandwidth

zmap 2.1.1 (custom build w/ JSON extension)

python 3.5.2

kafka-python 1.3.3

nginx 1.10

Scanning the entire public Internet is a political and technical challenge.

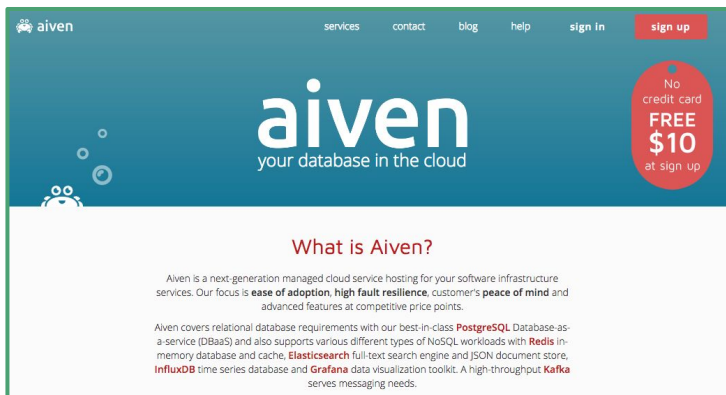
Public Notification

- Custom website w/ SSL cert
- DNS records
- Whois information
- Ability to opt-out

Technical Elements

- High-performance network I/O on AWS
- Zmap network scanner
- Scan management with Python
- Data streaming with Kafka

Kafka-as-a-Service: Aiven



The Aiven homepage features a dark blue header with the Aiven logo and navigation links: services, contact, blog, help, sign in, and sign up. The main content area has a large 'aiven' logo with the tagline 'your database in the cloud'. A red badge on the right says 'No credit card FREE \$10 at sign up'. Below the logo, the text 'What is Aiven?' is followed by a paragraph describing Aiven as a next-generation managed cloud service. A list of supported services is provided: PostgreSQL, Redis, Elasticsearch, InfluxDB, and Grafana.

What is Aiven?

Aiven is a next-generation managed cloud service hosting for your software infrastructure services. Our focus is **ease of adoption**, **high fault resilience**, customer's **peace of mind** and advanced features at competitive price points.

Aiven covers relational database requirements with our best-in-class **PostgreSQL** Database as-a-service (DBaaS) and also supports various different types of NoSQL workloads with **Redis** in-memory database and cache, **Elasticsearch** full-text search engine and JSON document store, **InfluxDB** time series database and **Grafana** data visualization toolkit. A high-throughput **Kafka** serves messaging needs.

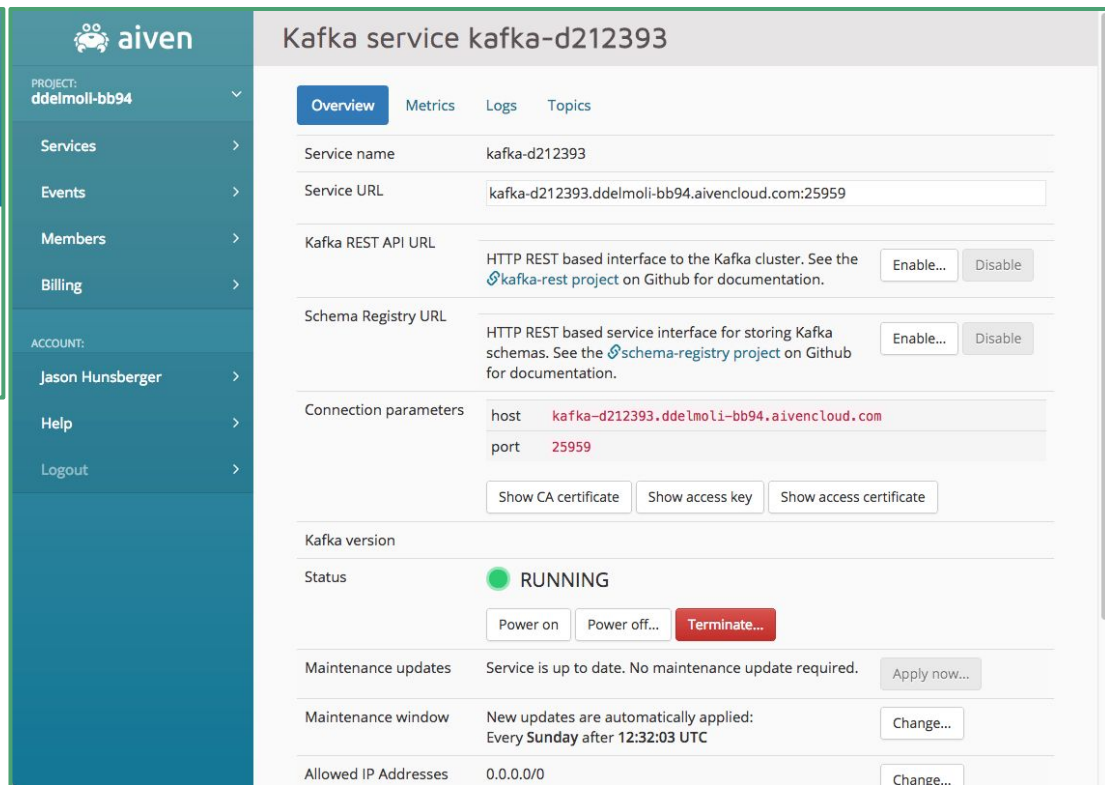
Kafka Cluster

aiven.io Kafka-as-a-Service

kafka 0.10.2

Startup plan 3 node HA set

Secured via client certificates and SSL



The Aiven Kafka service dashboard shows the configuration for the 'kafka-d212393' service. The left sidebar contains navigation links for Project, Services, Events, Members, Billing, Account, and Help. The main content area displays the service details, including the service name, URL, REST API URL, Schema Registry URL, connection parameters, Kafka version, status, maintenance updates, maintenance window, and allowed IP addresses.

Kafka service kafka-d212393

Overview Metrics Logs Topics

Service name: kafka-d212393

Service URL: kafka-d212393.ddelmoli-bb94.aivencloud.com:25959

Kafka REST API URL: HTTP REST based interface to the Kafka cluster. See the [kafka-rest project](#) on Github for documentation. Enable... Disable

Schema Registry URL: HTTP REST based service interface for storing Kafka schemas. See the [schema-registry project](#) on Github for documentation. Enable... Disable

Connection parameters

host: kafka-d212393.ddelmoli-bb94.aivencloud.com

port: 25959

Show CA certificate Show access key Show access certificate

Kafka version

Status: ● **RUNNING**

Power on Power off... Terminate...

Maintenance updates: Service is up to date. No maintenance update required. Apply now...

Maintenance window: New updates are automatically applied: Every Sunday after 12:32:03 UTC Change...

Allowed IP Addresses: 0.0.0.0/0 Change...


Apache Storm Processing Implementation details

Storm DAG

- Receive Kafka stream
- Convert IP to decimal
- Enrich with Country codes
- Scan-level analytics

Storm Processor

AWS EC2 Amazon Linux
python 2.7.12 (compiled with SSL)
kafka-python 1.3.3
java 1.7
apache storm 0.9.3
streamparse 2.1.4

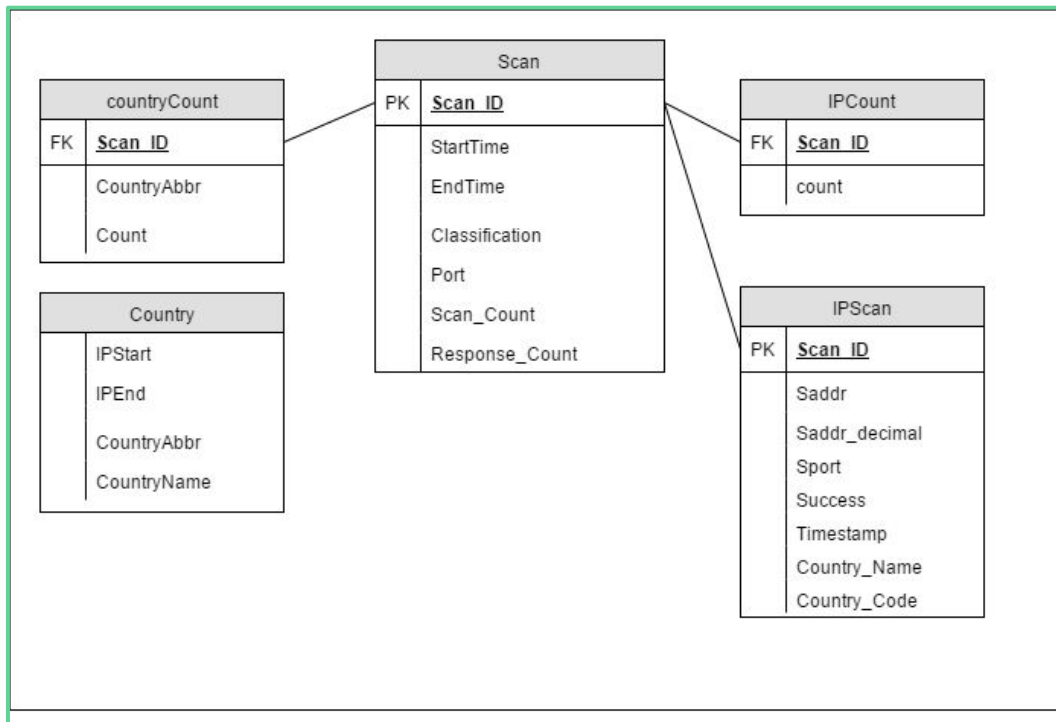


Postgres Database details

Database

Amazon AWS RDS PostgreSQL

1 TB of storage



Technical Challenges

Front End

Challenge: needed access to high-performance hardware

Solution: use AWS HVM

Challenge: needed to manage long-running tools and create repeatable, dependable output

Solution: Python wrapper classes

Challenge: needed to supplement Zmap-generated output

Solution: use JSON output, purchase Country to IP mapping DB, and augment JSON records per scan.

Storm Processing

Challenge: Kafka cluster connections require SSL and python $\geq 2.7.9$

W205 AMIs use python 2.7.3 and lack SSL

Unable to add SSL / upgrade python on W205 AMIs

Solution: build new AMI with python 2.7.12 and SSL

Challenge: latest versions of apache storm (1.1.0) & streamparse (3.4.0) are incompatible; latest version of streamparse (3.4.0) uses non-Clojure topology definitions

Solution: Install apache storm 0.9.3 and streamparse 2.1.4 to mimic successful W205 AMI installation

Project Results

As of today, we are able to:

- Scan the Internet with both TCP and ICMP scans
- Send data into our Kafka-as-a-Service cluster
- Store/lookup data in Postgres

What we have left:

- Confirm our analytics are correct in Storm
- Scale the infrastructure to meet our processing needs

```
'target_port': 80}
>>> icmp_scanner = IcmpScanner('zmap-scan.conf', max=4e6, output_file='/data/icmp-scan001.json')
>>> pprint(icmp_scanner.settings())
{'config_file': 'zmap-scan.conf',
 'cooldown_time': 9,
 'max_targets': 4000000,
 'output_fields': ['saddr',
                  'daddr',
                  'ttl',
                  'classification',
                  'success',
                  'cooldown',
                  'timestamp_str'],
 'output_file': '/data/icmp-scan001.json',
 'probe_module': 'icmp_echo_scan',
 'target_port': 0}
>>> http_scanner.run()
0:00 0%; send: 0 0 p/s (0 p/s avg); rcv: 0 0 p/s (0 p/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.00%
0:00 0%; send: 1290 459 Kp/s (30.4 Kp/s avg); rcv: 0 0 p/s (0 p/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.00%
0:01 7%; send: 746527 745 Kp/s (716 Kp/s avg); rcv: 906 905 p/s (868 p/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.12%
0:02 14%; send: 1454556 708 Kp/s (712 Kp/s avg); rcv: 3014 2.11 Kp/s (1.48 Kp/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.21%
0:03 21%; send: 2162033 707 Kp/s (710 Kp/s avg); rcv: 4942 1.93 Kp/s (1.62 Kp/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.23%
0:04 28%; send: 2880921 719 Kp/s (712 Kp/s avg); rcv: 6429 1.49 Kp/s (1.59 Kp/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.22%
0:05 34% (10s left); send: 3582197 701 Kp/s (710 Kp/s avg); rcv: 7943 1.51 Kp/s (1.57 Kp/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.22%
0:06 41% (9s left); send: 4000000 done (707 Kp/s avg); rcv: 9401 1.46 Kp/s (1.55 Kp/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.24%
0:07 48% (8s left); send: 4000000 done (707 Kp/s avg); rcv: 9992 590 p/s (1.42 Kp/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.25%
0:08 55% (7s left); send: 4000000 done (707 Kp/s avg); rcv: 10001 8 p/s (1.24 Kp/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.25%
0:09 62% (6s left); send: 4000000 done (707 Kp/s avg); rcv: 10013 11 p/s (1.11 Kp/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.25%
0:10 69% (5s left); send: 4000000 done (707 Kp/s avg); rcv: 10020 6 p/s (997 p/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.25%
0:11 75% (4s left); send: 4000000 done (707 Kp/s avg); rcv: 10026 5 p/s (907 p/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.25%
0:12 82% (3s left); send: 4000000 done (707 Kp/s avg); rcv: 10027 0 p/s (832 p/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.25%
0:13 89% (2s left); send: 4000000 done (707 Kp/s avg); rcv: 10032 4 p/s (768 p/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.25%
0:14 96% (1s left); send: 4000000 done (707 Kp/s avg); rcv: 10032 0 p/s (714 p/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.25%
>>> icmp_scanner.run()
0:00 0%; send: 2216 0 p/s (61.0 Kp/s avg); rcv: 0 0 p/s (0 p/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.00%
0:01 7%; send: 701802 698 Kp/s (676 Kp/s avg); rcv: 4367 4.36 Kp/s (4.21 Kp/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.62%
0:02 14%; send: 1511960 810 Kp/s (742 Kp/s avg); rcv: 5286 918 p/s (2.59 Kp/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.35%
0:03 21%; send: 2332440 820 Kp/s (768 Kp/s avg); rcv: 5352 65 p/s (1.76 Kp/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.23%
0:04 29%; send: 3151471 819 Kp/s (780 Kp/s avg); rcv: 5363 10 p/s (1.33 Kp/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.17%
0:05 36% (10s left); send: 3981448 830 Kp/s (790 Kp/s avg); rcv: 5372 8 p/s (1.07 Kp/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.13%
0:06 43% (9s left); send: 4000000 done (790 Kp/s avg); rcv: 5376 3 p/s (890 p/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.13%
0:07 50% (8s left); send: 4000000 done (790 Kp/s avg); rcv: 5379 2 p/s (764 p/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.13%
0:08 57% (7s left); send: 4000000 done (790 Kp/s avg); rcv: 5380 0 p/s (669 p/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.13%
0:09 64% (6s left); send: 4000000 done (790 Kp/s avg); rcv: 5382 1 p/s (595 p/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.13%
0:10 71% (5s left); send: 4000000 done (790 Kp/s avg); rcv: 5382 0 p/s (536 p/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.13%
0:11 79% (4s left); send: 4000000 done (790 Kp/s avg); rcv: 5382 0 p/s (487 p/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.13%
0:12 86% (3s left); send: 4000000 done (790 Kp/s avg); rcv: 5382 0 p/s (446 p/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.13%
0:13 93% (2s left); send: 4000000 done (790 Kp/s avg); rcv: 5383 0 p/s (412 p/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.13%
0:14 100% (1s left); send: 4000000 done (790 Kp/s avg); rcv: 5383 0 p/s (383 p/s avg); drops: 0 p/s (0 p/s avg); hitrate: 0.13%
```

Limitations and Future Extensions

- Limitations

- The cost of the infrastructure was beginning to grow by the end of the project
- The front end would scan even faster using PF_RING network drivers.
- Later versions of Storm & Streamparse may be more efficient

- Future Extensions / Additions

- Expand beyond TCP and ICMP scans
- Enable different JSON record formats depending upon the type of scan
- Provide not only Country-level data, but also city-level and ISP-level lookups
- Sharded or more scalable database design for analysis (i.e., AWS Aurora PostgreSQL or AWS RedShift)
- Hosted query option for data
- Geographic distribution of scanning servers
- HTML visualizations for our streaming and batch analytics

Thank you

Appendix